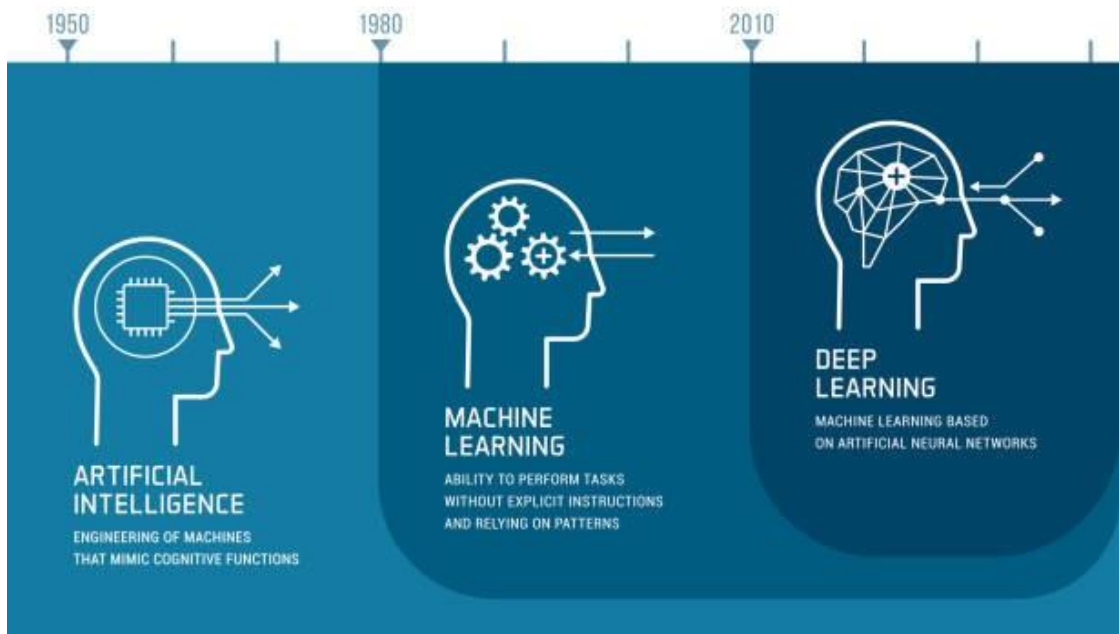


Intro to Neural Networks

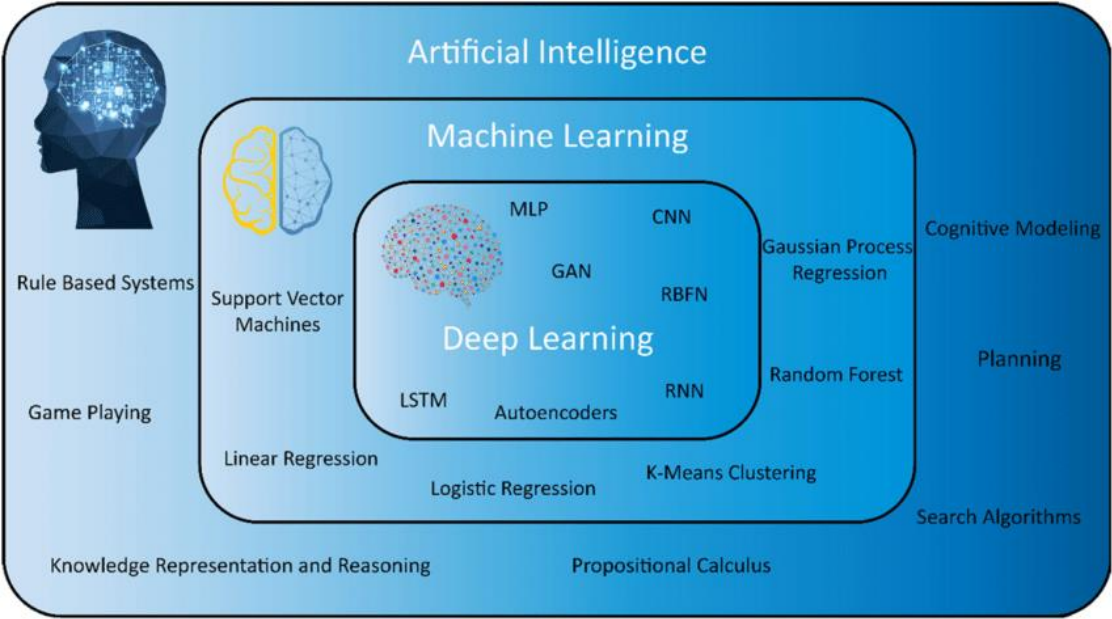
BA865 – Mohannad Elhamod

Motivation

What is Deep Learning?





[LinkedIn](#)



[ResearchGate](#)

Behold The Almighty Deep Learning!




<https://thispersondoesnotexist.com>



Behold The Almighty Deep Learning!


- [Generate music](#)
- [Classify images](#)
- [Classify images \(2\)](#)
- [Dimensionality Reduction](#)

Beware The Hype!



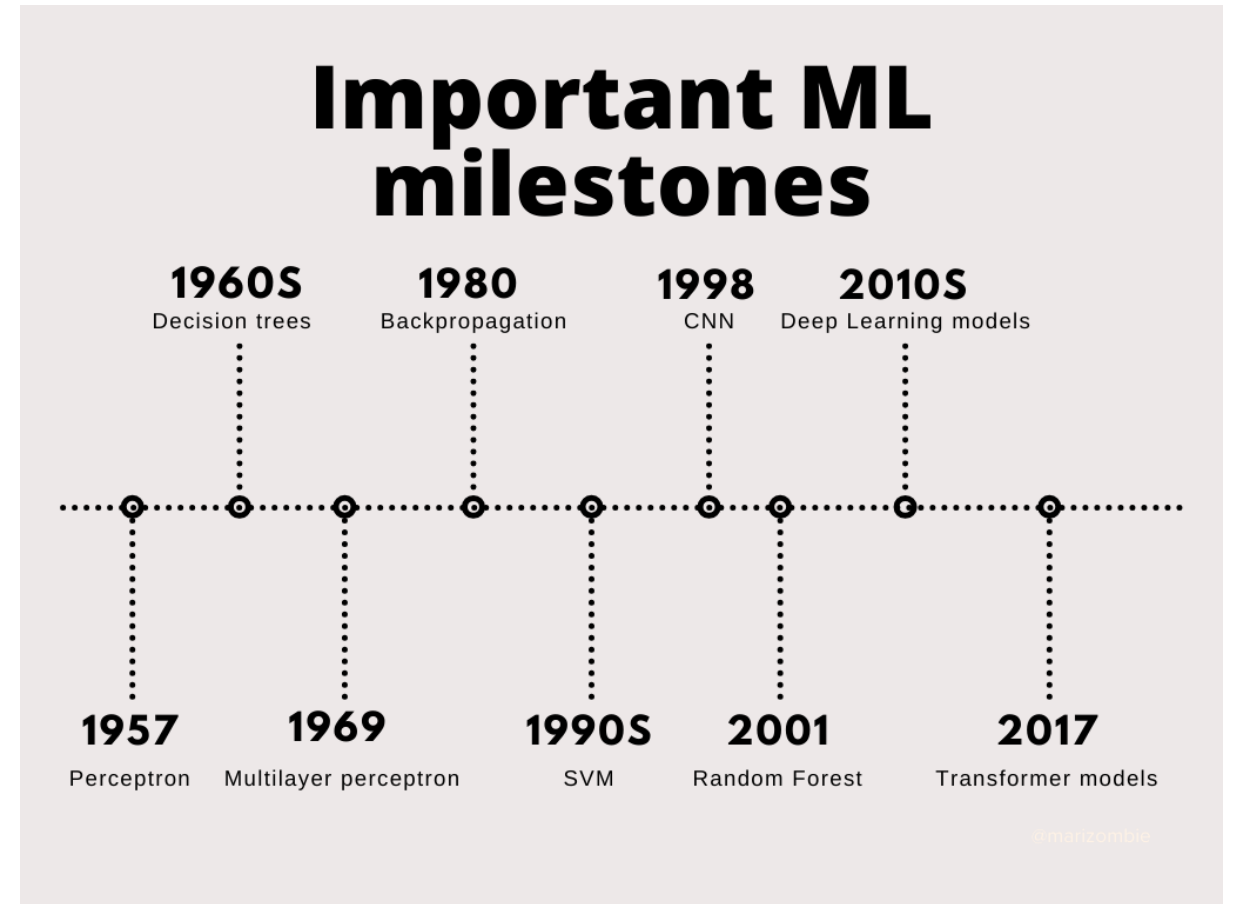
Beware The Hype!

- It hasn't always looked rosy...There are some "hypes" along the way.



Algorithms Come and Go...


- Different algorithms gain popularity at different periods of time.



medium.com

Why Deep Learning?

- Traditional ML was found to plateau in performance in the “Big Data” realm.
- Deep Learning has shown a much better capacity for learning from large amounts of data.



Why Now?


Geoffrey Hinton, Godfather of AI.

What was actually wrong with backpropagation in 1986?

- We all drew the wrong conclusions about why it failed. The real reasons were:

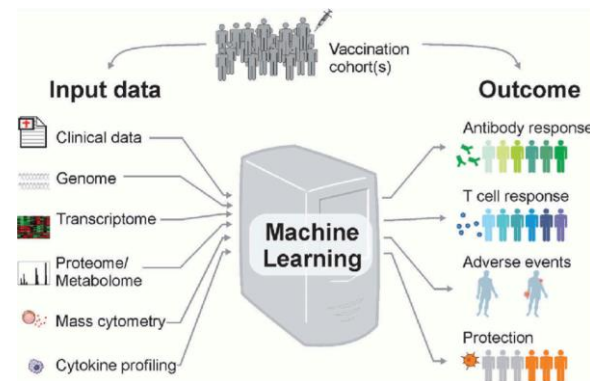

1. Our labeled datasets were thousands of times too small.
2. Our computers were millions of times too slow.
3. We initialized the weights in a stupid way.
4. We used the wrong type of non-linearity.

THE ROYAL SOCIETY

A photograph of Geoffrey Hinton, a man with grey hair wearing a dark suit and white shirt, standing at a wooden podium. He is speaking into a microphone. Behind him is a blue backdrop with white text that reads "Backpropagation: science, technology, and the future". To his left, on the podium, are two clear water bottles. The background of the slide is black with the words "THE ROYAL SOCIETY" in large, white, serif capital letters.

Where Can Deep Learning Be Applied?

- Almost anywhere where you have data to learn from and make decisions based on.

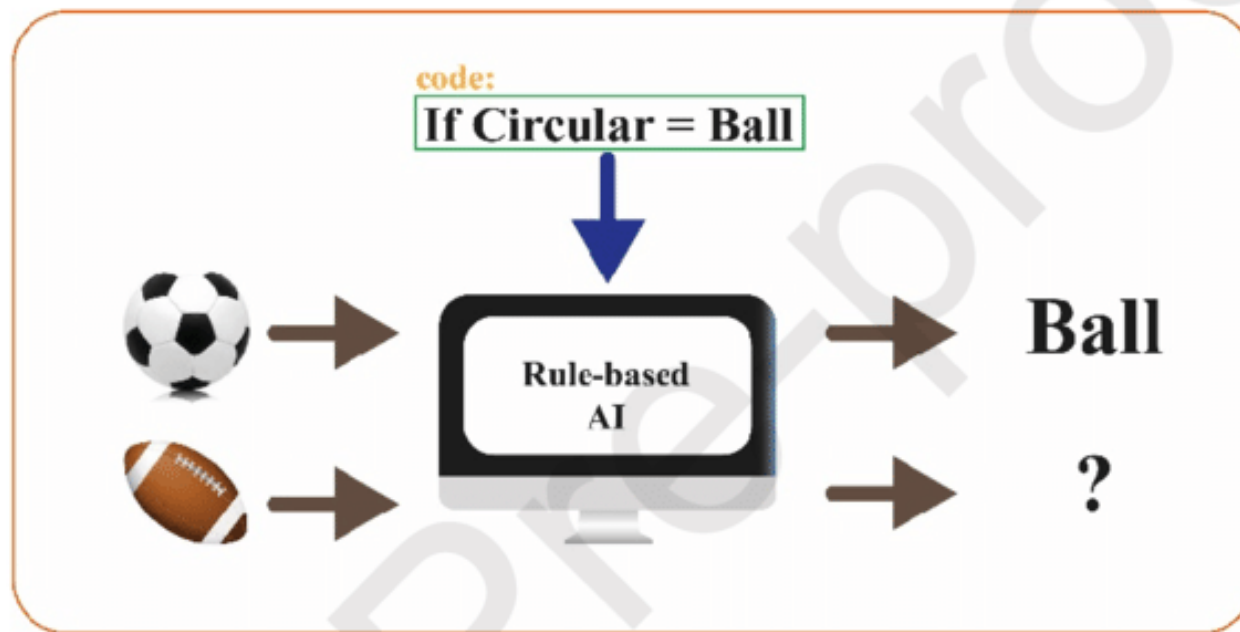


Machine Learning:

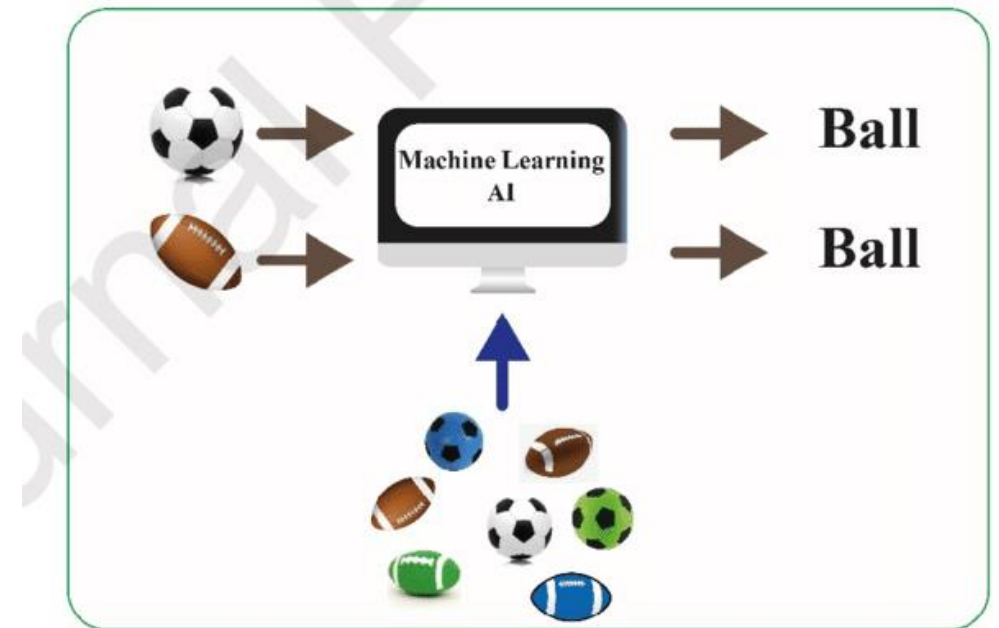
An Essential Review!

What is Machine Learning?

Elbadawi, Moe & McCoubrey, Laura & Gavins, Francesca & Ong, Jun Jie & Goyanes, Alvaro & Gaisford, Simon & Basit, Abdul. (2021). Harnessing Artificial Intelligence for the Next Generation of 3D Printed Medicines. *Advanced Drug Delivery Reviews*. 175. 10.1016/j.addr.2021.05.015.




Knowledge-based modeling



Data-driven modeling


What is Machine Learning?

<https://www.wangzze.top>



- Explicit programming is used to solve problems.
- Rules can be manually specified.

Knowledge-based modeling

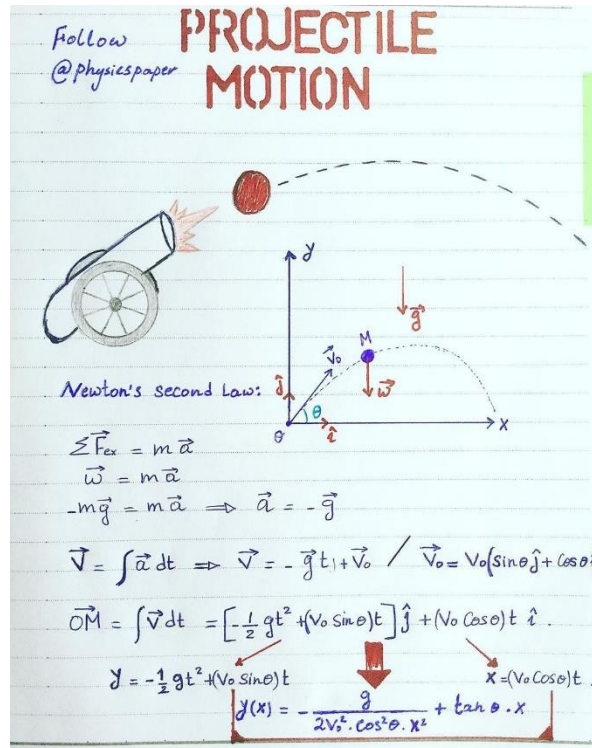


- Samples are used for training.
- The decision-making rules are complex or difficult to describe.
- Rules are automatically learned by machines.

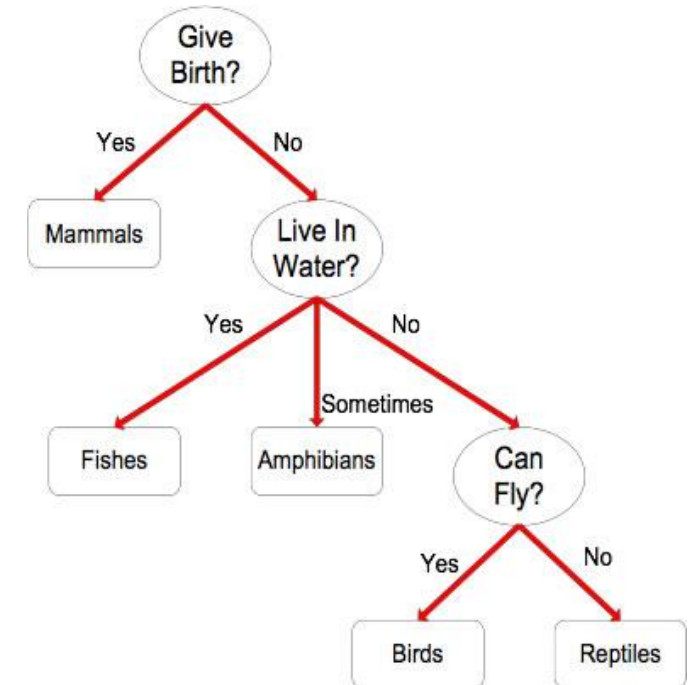
Data-driven modeling

Knowledge-based Modeling


- What challenges you find with this approach?




<http://jcsites.juniata.edu/faculty/rhodes/ml/rulebasedClass.htm>
<https://www.pinterest.com/pin/598626975454754006/>



Supervised Learning




MNIST Dataset




Phase 1: Training

Supervised Learning

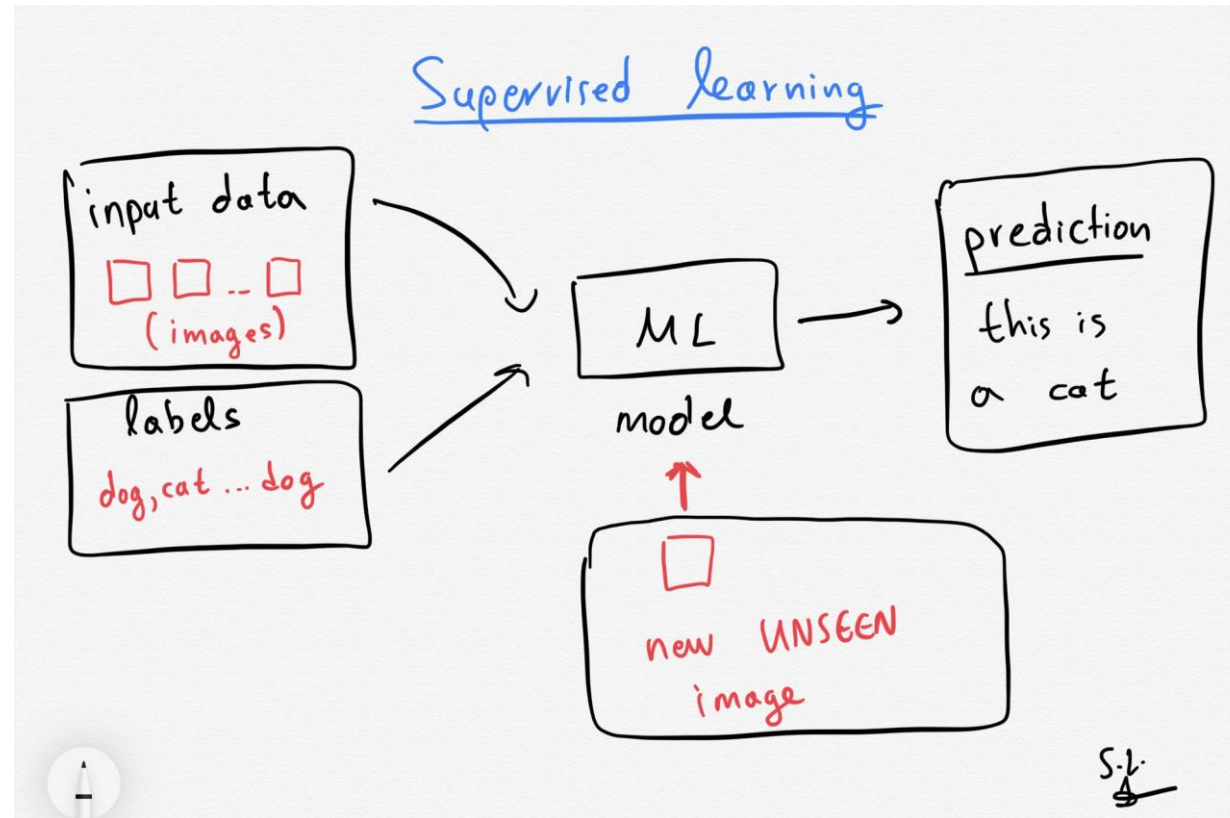


MNIST Dataset




Phase 2: Validation/Testing

Supervised Learning



<https://towardsdatascience.com/what-is-machine-learning-a-short-note-on-supervised-unsupervised-semi-supervised-and-aed1573ae9bb>

Supervised Learning - Model Types



<https://aldro61.github.io/microbiome-summer-school-2017/sections/basics/>

Supervised Learning - Model Types




Regression

What is the temperature going to be tomorrow?



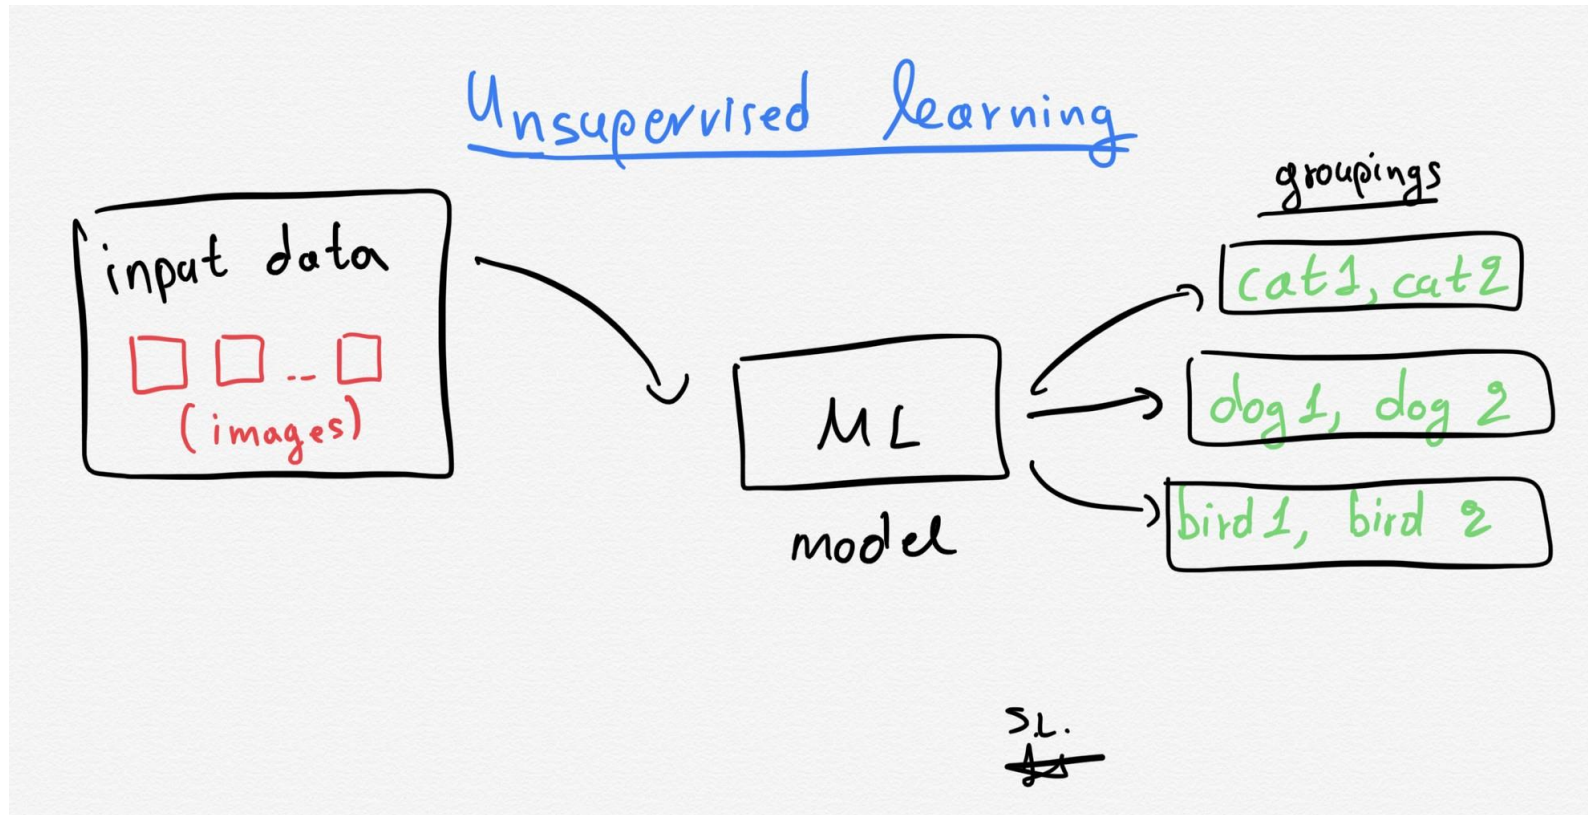
Classification

Will it be Cold or Hot tomorrow?




https://medium.com/@ali_88273/regression-vs-classification-87c224350d69

Unsupervised Learning




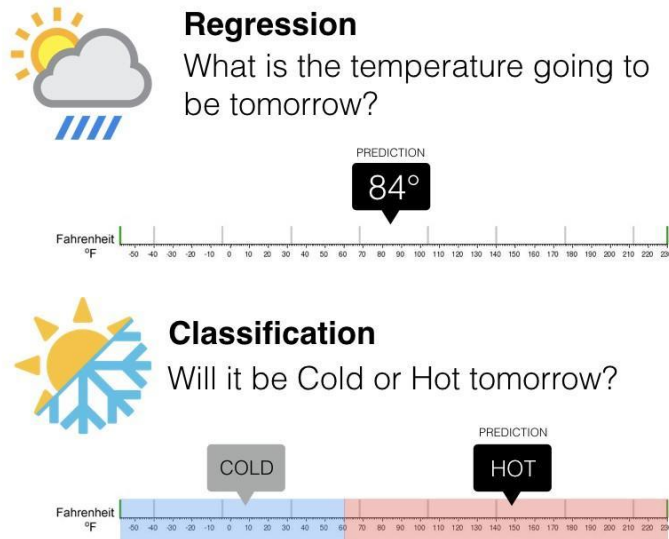
<https://towardsdatascience.com/what-is-machine-learning-a-short-note-on-supervised-unsupervised-semi-supervised-and-aed1573ae9bb>

The Pillars of ML



Input & Output

- What exactly are you trying to model?



<https://www.wired.com/2015/04/cost-of-sensors-autonomous-cars>
https://medium.com/@ali_88273/regression-vs-classification-87c224350d69

Data

- Is there enough of it?
- Does it need clean-up?



Damaged specimen




Missing Features





Occluded Features

Elhamod, M., Diamond, K. M., Murat Maga, A., Bakis, Y., Bart, H. L., Mabee, P., Dahdul, W., Leipzig, J., Greenberg, J., Avants, B., & Karpayne, A. (2022). Hierarchy-guided neural network for species classification. *Methods in Ecology and Evolution*, 13, 642–652. <https://doi.org/10.1111/2041-210X.13768>



Model

- What kind of model would be sufficient/suitable for modeling your data?




<https://statisticallyrelevant.com/decision-trees-in-python-predicting-diabetes/>

Resources


- What compute resources are available?
- What is the project's timeline?

Model Generalization




Underfitting

$$\theta_0 + \theta_1 x$$



Well-fitted

$$\theta_0 + \theta_1 x + \theta_2 x^2$$



Overfitting

$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

<https://datascience.foundation/sciencewhitepaper/underfitting-and-overfitting-in-machine-learning>

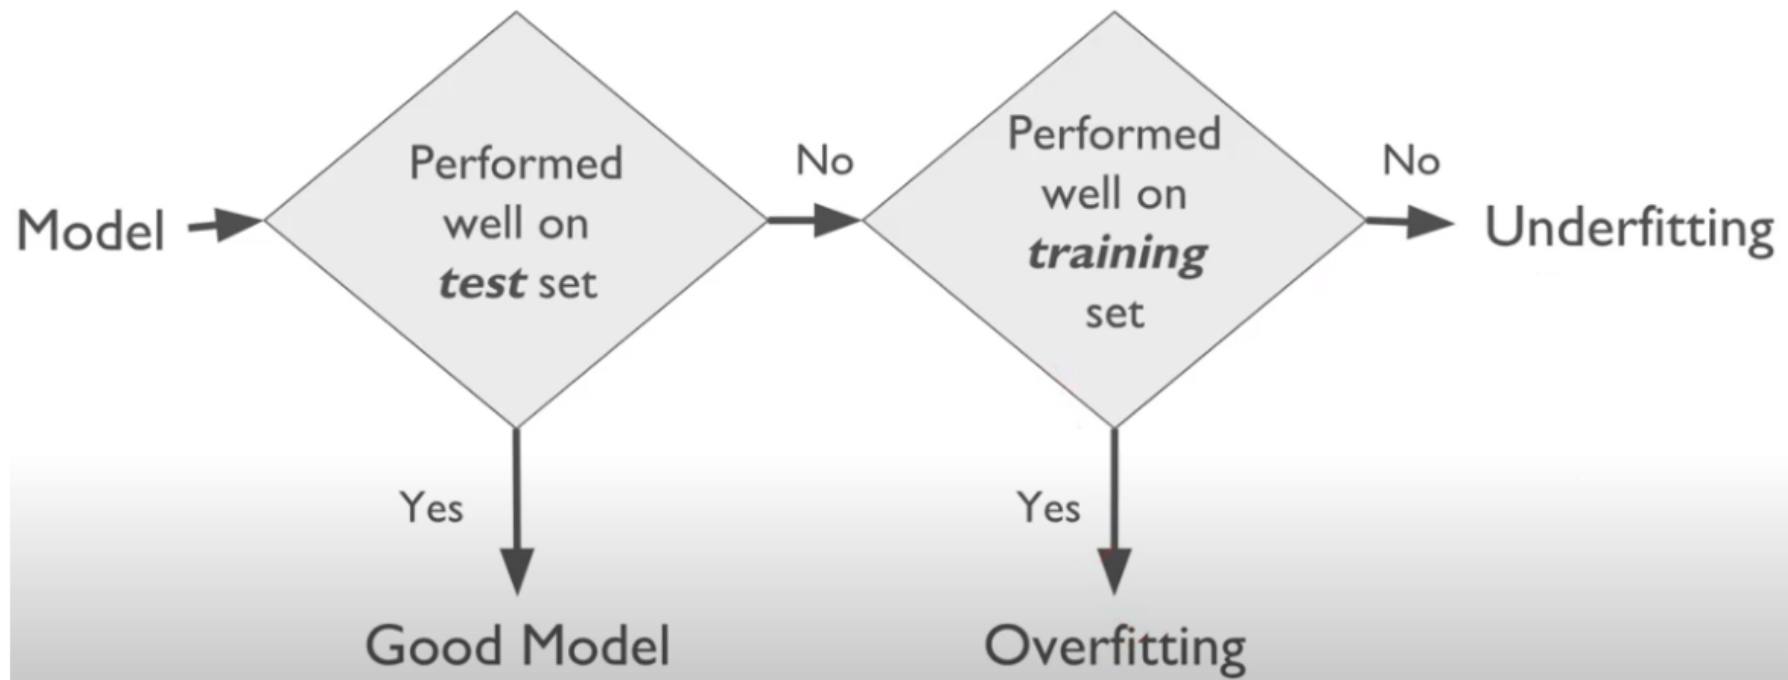
What is Model Complexity?

- It is the degrees of freedom, the “flexibility” of the model to fit the data.
- Bigger datasets or datasets with more attributes need more complex models, and vice versa.
- [Try it.](#)

Demo

Put Your Own Subtitle In This Box

Are We Overfitting or Underfitting?




So, What If We Overfit?!

- Without validation/testing, you could overfit
 - In other words, you could learn the wrong patterns in the data!



<https://medium.com/trusted-ai/explaining-ai-model-behaviour-with-ibm-watson-openscale-86515702c177>


Training vs. Test Errors



www.researchgate.net

Training vs Test vs.?

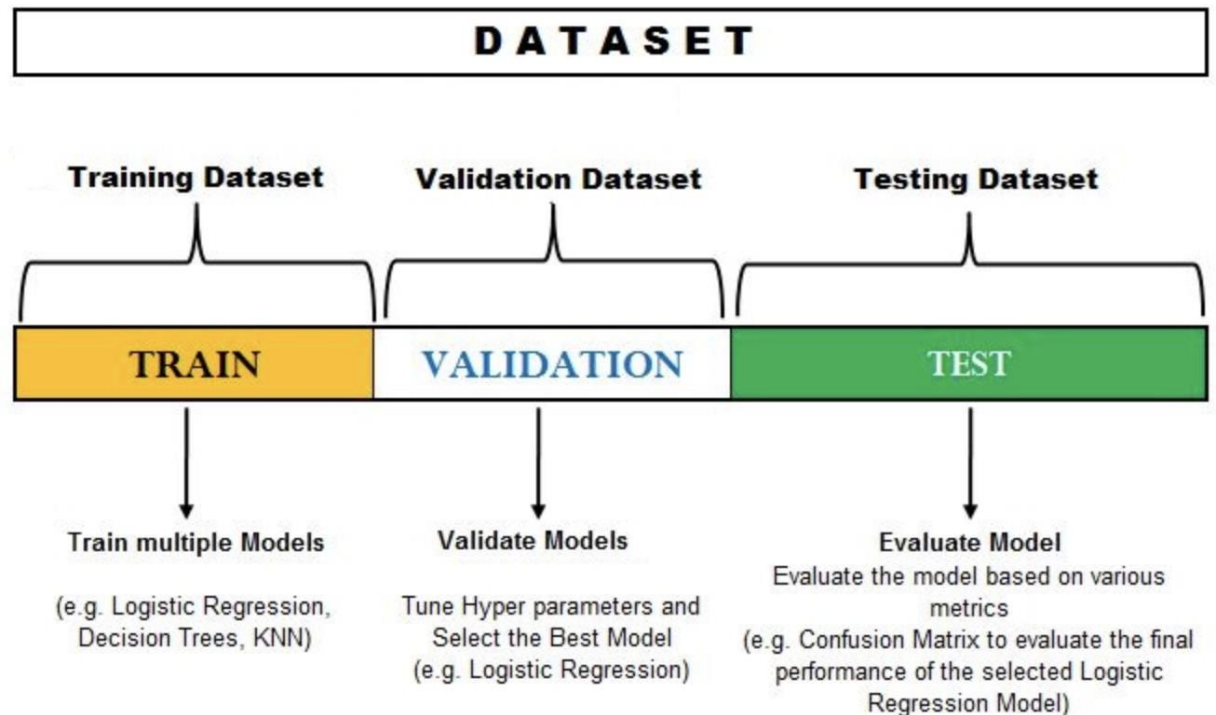
- How do we test the model's performance?
- How do we assess the model's fine-tuning?



<https://www.machinecurve.com/index.php/2020/11/16/how-to-easily-create-a-train-test-split-for-your-machine-learning-model/>

Training vs Test vs. Validation?

- The test set should only be used once our model is ready, it should not be used to “tune” the model.
- Hence, we need a third “validation” dataset, which is a split of the training set.



<https://vitalflux.com/hold-out-method-for-training-machine-learning-model/>

Logistics

Put Your Own Subtitle In This Box

Who is your professor?



Why are you here?

- I am sure it eventually leads to money....
- Not a required class. So, hopefully based on interest!
- Nonetheless, understanding how things work is essential for decision making and innovation.



Someone who had begun to [study] geometry asked Euclid, 'What shall I get by learning these things?' Euclid called his slave and said, 'Give him [some money], since he must make gain out of what he learns'.

(Heath, 1981, loc. 8625)



Euclid

Fundamentals are important!

- The more foundational knowledge you skip, the more fundamental errors you will make.
- Work hard. **Be patient!**



Zen Speaks: Shouts of Nothingness


Your professor is not a God

- Deep Learning is a fast-growing domain.
- The internet nowadays has all kinds of learning material.
- Your professor is **NOT** here as a walking encyclopedia. He is here to guide your learning experience and build you a solid foundation, so you could continue learning on your own later.




No Question is Foolish

- Other might have the same question.
- Knowledge is hierarchical.



Can I get an A?

- [Questrom policy](#)
- But you will get your fair chance.



Participation is Essential!

- 7.5% of your grade!
- The instructor reserves the right to cold-call.



Office Hours

- They are for you to take advantage of!
- However, to make the best of your and the TA's/instructor's time, do your homework before dropping in:
 - If you have a question about your project, make sure you have synched with your colleague in advance.
 - If you have question about assignments, make sure you have done your best and that your question is specific rather than “How do I solve this?”.



Syllabus

- [Here](#)
 - Pay attention to [Piazza](#)!
 - Attendance and participation.
 - The class is somewhat more technical than BA820 and BA810.
 - Zero tolerance for cheating (AI or peers).
 - 5%-20% Penalties for deviating from assessment instructions!

Shared Computing Cluster (SCC)

- BU's computing cluster with access to CPUs and GPUs.
- We have a limited number of compute hours. Use your allocations wisely!
- When facing issues, ask the professor or TA for help first, not IT support!
- **Don't leave things till last minute!**



Coding Environments

- <https://www.simplilearn.com/keras-vs-tensorflow-vs-pytorch-article>
 - Keras is easier to use for non-CS users. It is good at an introductory level.
 - PyTorch is most commonly used for research as it provides powerful low-level API.
- We will also use some more advanced packages that help us diagnose and use more advanced models.



Hugging Face



Weights & Biases