

# Intro to Neural Networks

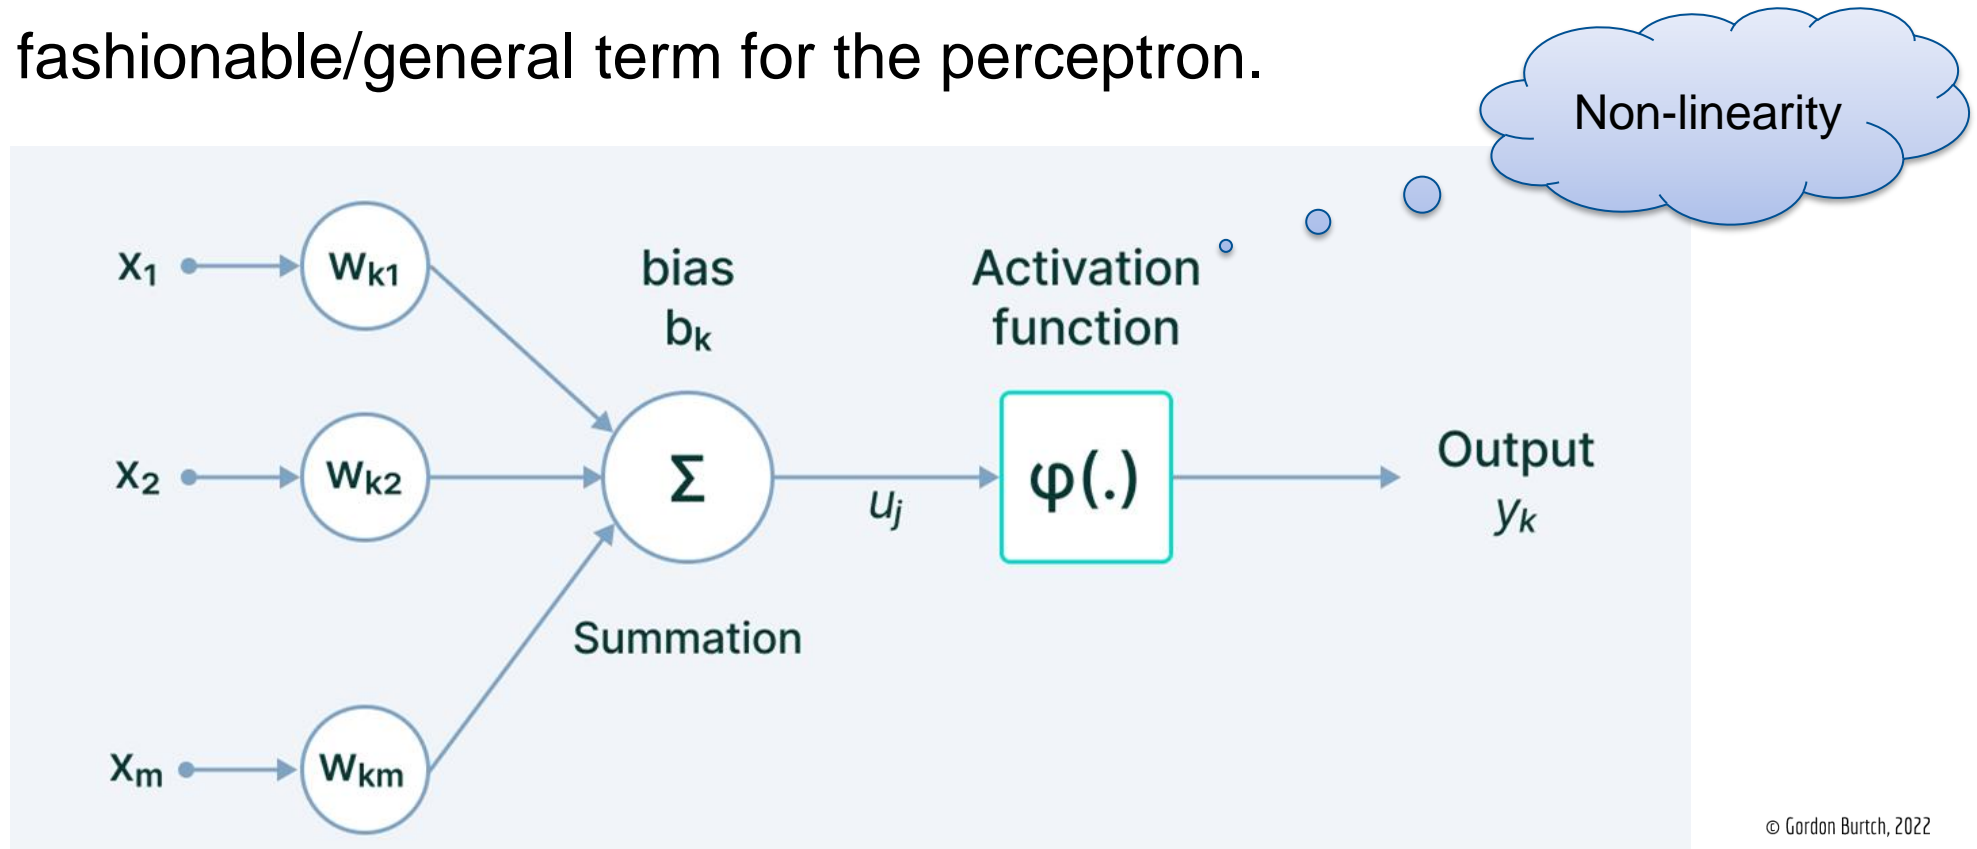
**BA865 – Mohannad Elhamod**

# MLPs

## The Multi-Layer Perceptron

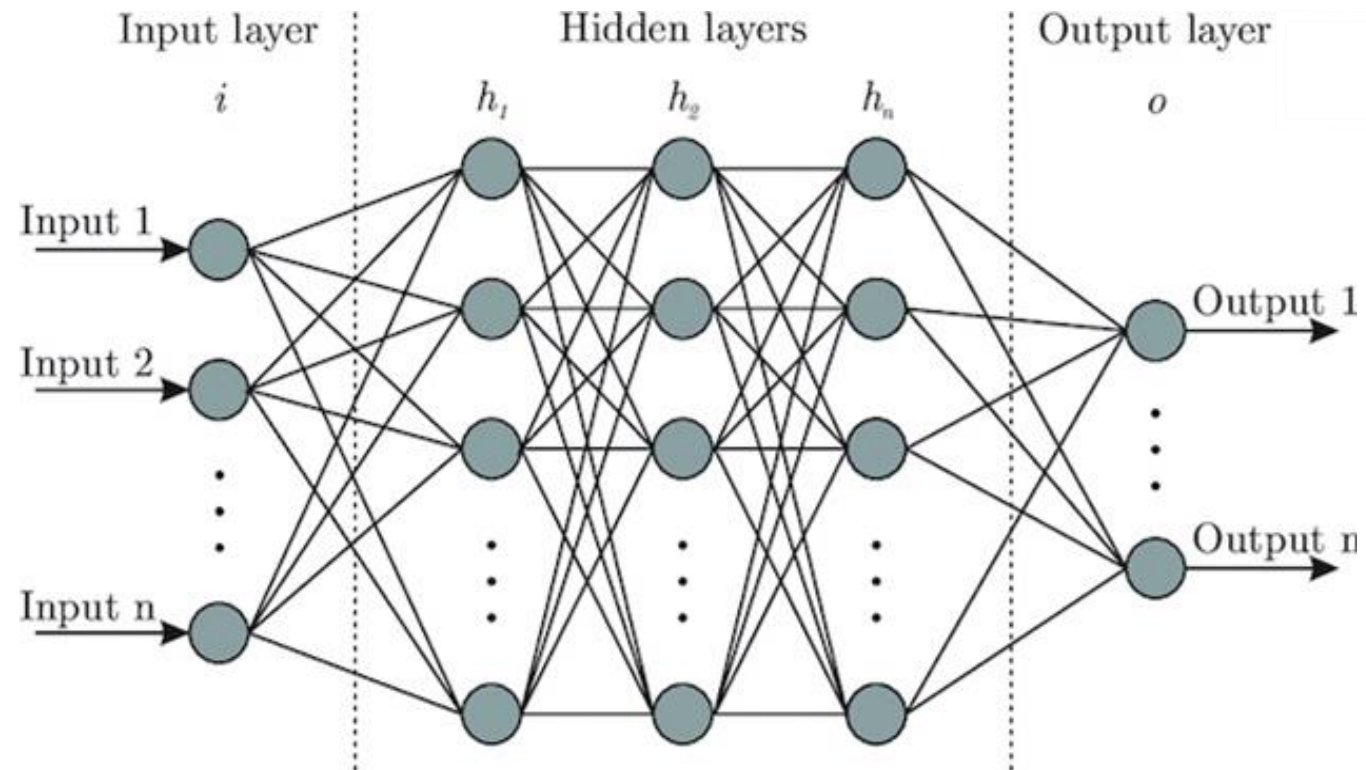
# A Neuron

- A more fashionable/general term for the perceptron.



© Gordon Burtch, 2022

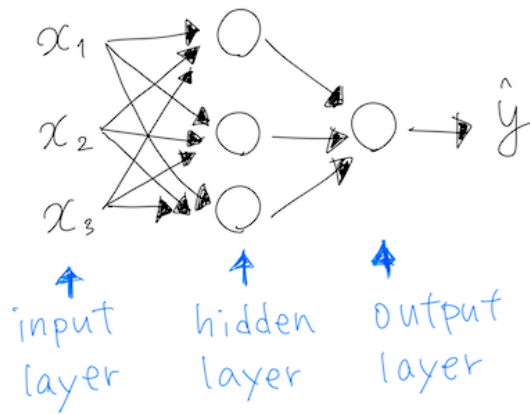
# Neural Networks



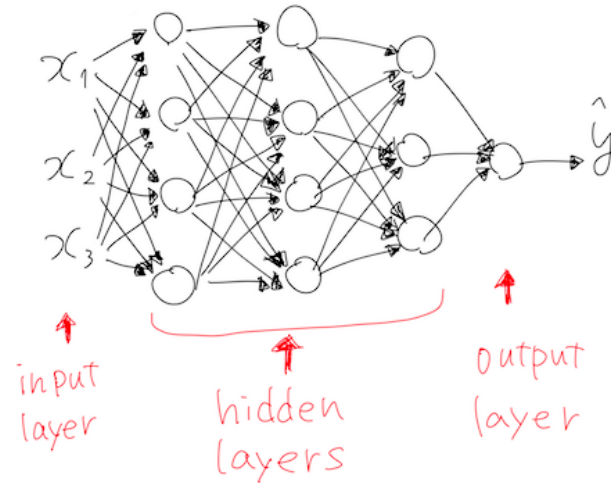
# Deep Networks

- Deep = More and more layers...
- leading to more complexity and better capacity for capturing complex phenomena.

Shallow Neural Network




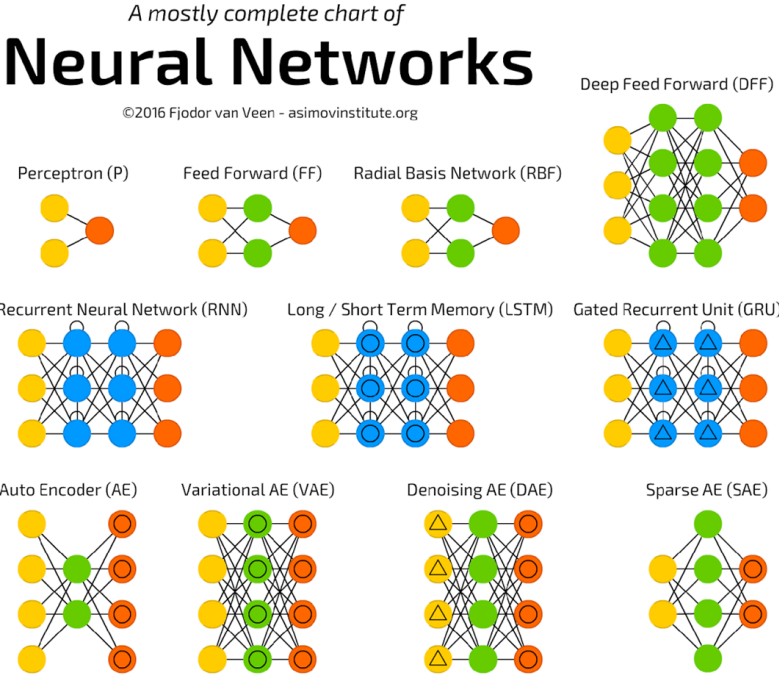
Deep Neural Network



# MLPs, One of Many Types...

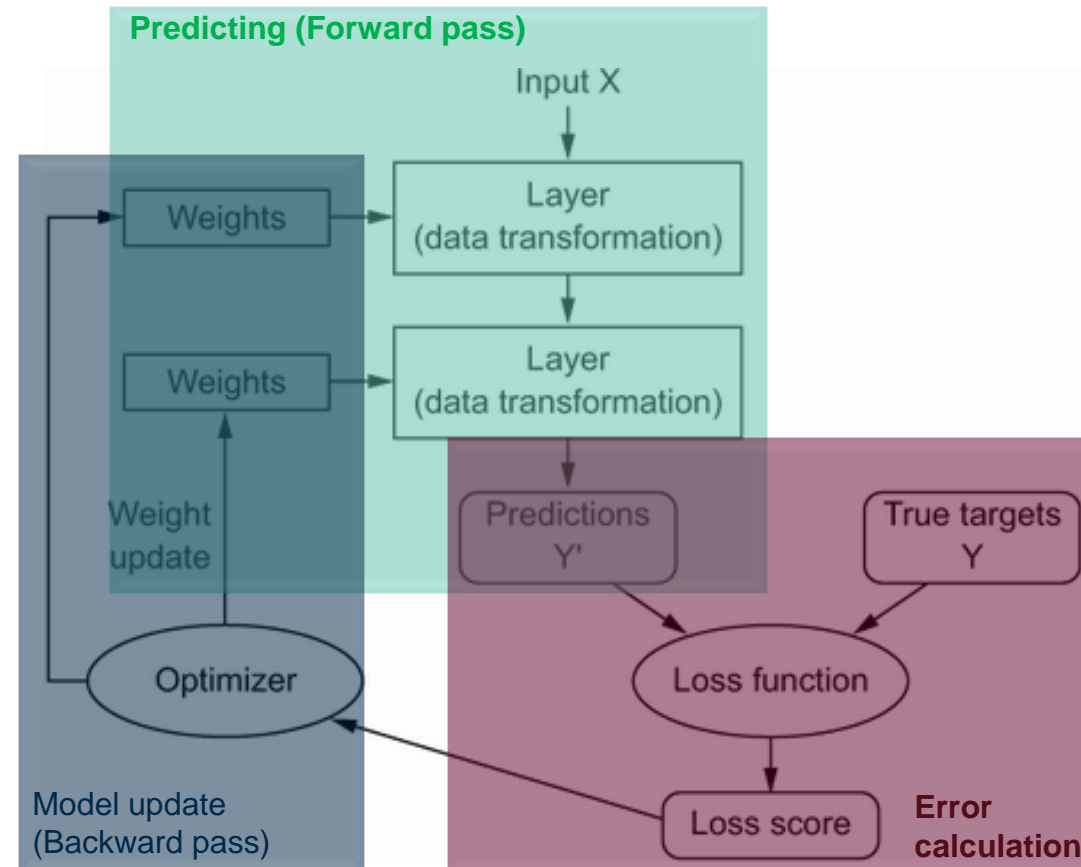
- MLP = FF (Feed Forward) Network.

- Backfed Input Cell
- Input Cell
- Noisy Input Cell
- Hidden Cell
- Probablistic Hidden Cell
- Spiking Hidden Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- Different Memory Cell
- Kernel
- Convolution or Pool



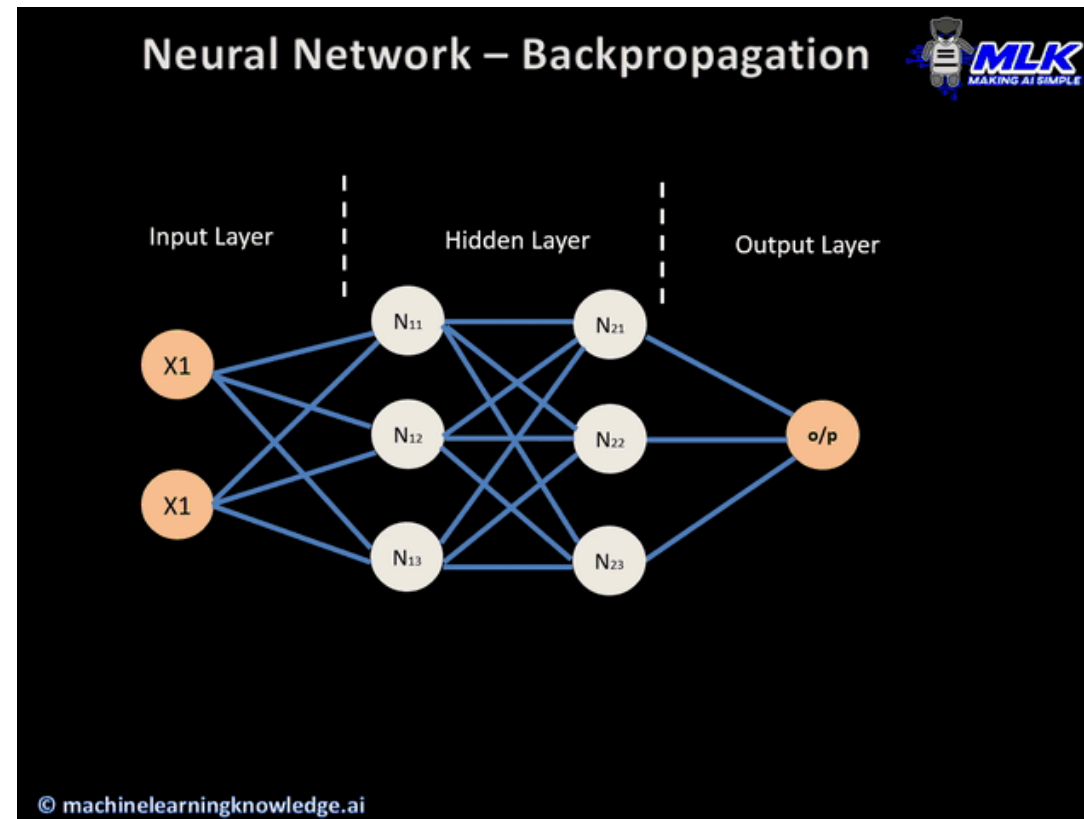
# Disecting The Neural Network

# The Framework



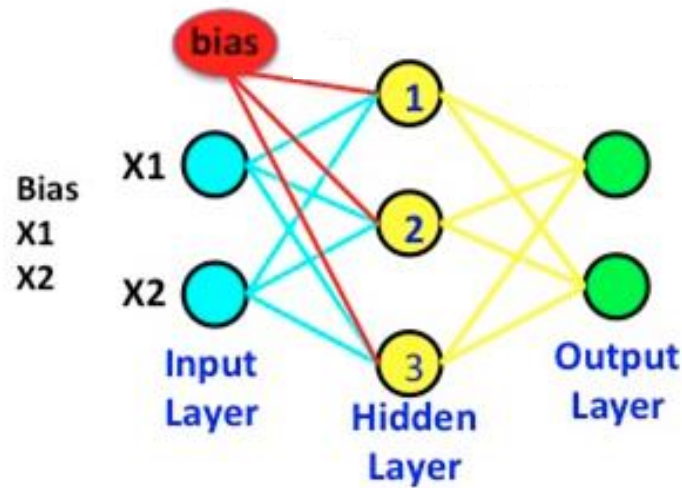


# The Framework



# Predicting

- What is a layer actually doing?
- Each layer is a matrix multiplication followed by a non-linearity!
  - Why bother with the non-linearity?!



Input Layer

bias X1 X2

$$\begin{bmatrix} 1 & 0 & 1 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} .5 & .5 & .5 \\ .5 & .5 & .5 \\ .5 & .5 & .5 \end{bmatrix} =$$

Weights  $w^T$  (transposed)

Go to Hidden Nodes

Hidden Layer

1 2 3

3 x 3

Hidden Layer

Bias Node 1 Node 2 Node 3

$$\begin{bmatrix} 1 & 1 & 1 \\ .5 & .5 & .5 \\ .5 & .5 & .5 \\ 1 & 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} .2 & .1 \\ .4 & .1 \\ .4 & .1 \end{bmatrix} =$$

Weights

3 x 2

Output Layer

1 .3 .5 .15 .5 .15 1 .3

4 x 2

Sigmoid Function

$$\frac{1}{1 + e^{-(wx+b)}}$$

Figure (modified) courtesy of Baber Zaimen

# Predicting

- [Demo](#)
- The non-linearities allow the neural net to “warp” a non-linear problem into a linear one!

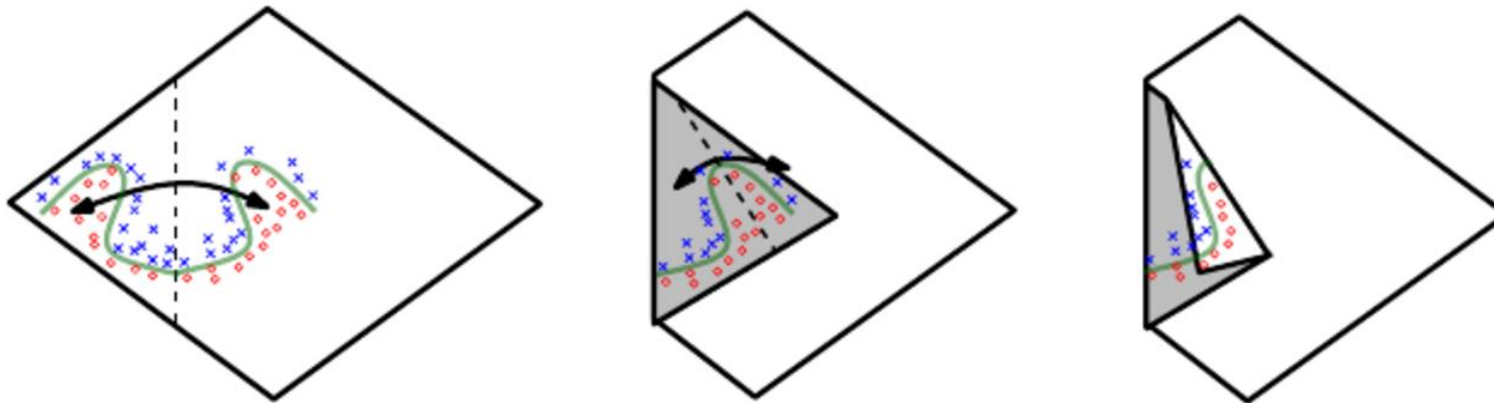


Figure courtesy of [Deep Learning Book](#)

# Optimization

- Using Gradient Descent (or some other optimizer) to “update the network.”
  - What do we exactly mean by “updating the network”?

# Optimization

- Gradient descent is performed with respect to the weights/biases.
- Behold...

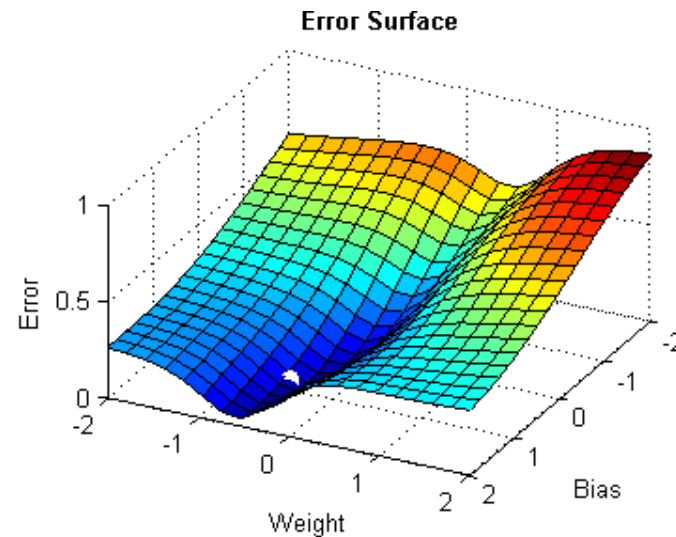
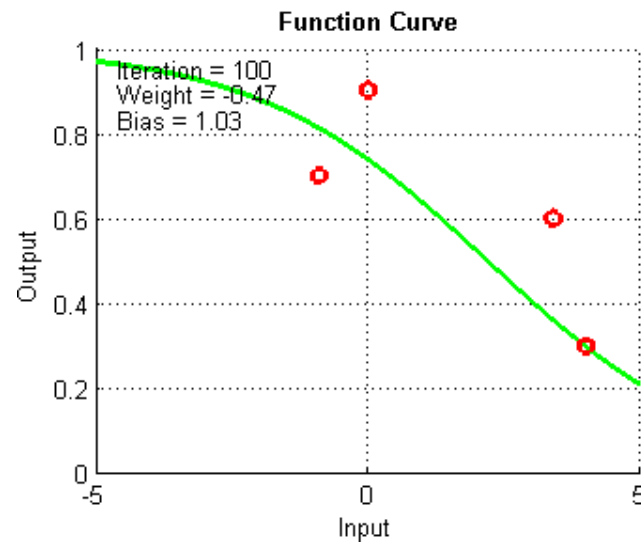



Figure courtesy of [Devin Soni](#)



# Optimization

- But how does the update get carried all the way back?
  - Chain rule!
  - This is called back-propagation.
- Luckily, these calculations can be automated with automatic differentiation.



# Optimization

- We need to make sure the gradient is non-zero...
  - Otherwise, the gradient can't "flow"!
- Replace the step function with a continuous one!



# Hyper-Parameters



# Learning Rate

- Generally, the most important hyperparameter of them all!
  - Too low: Really slow convergence.
  - Too high: No convergence.




Figure courtesy of [Stanford CS class CS231n](#)

# Learning Rate: Schedulers

- To get the best of both worlds, you could adjust the learning rate in phases.
  - This way, you still converge but faster.
- Using a scheduler is a common practice.




Figure courtesy of [B. D. Hammel](#)

# Learning Rate: Early Stopping

- The number of epochs impacts the model's fitness.
- Training needs to stop at the “right” epoch.
- How do we achieve that?
  - Better to stop when validation error stops decreasing for a certain number (n) of epochs.
  - Setting n too small or too large will impact convergence.




Figure courtesy of RAHUL JAIN

# Optimizer: Momentum

- Adding a momentum term (i.e., gradients from previous epochs), helps the convergence process.




Figure courtesy of [Fred Park](http://fredpark.com)





Figure courtesy of [Casper Hansen](#)

$$\begin{aligned} z^{k+1} &= \beta z^k + \nabla f(w^k) \\ w^{k+1} &= w^k - \alpha z^{k+1} \end{aligned}$$

Momentum coefficient

# Optimization: Batches

- Datasets are usually huge and won't fit in GPU memory in its entirety.
- So, we split the dataset into batches.
  - This is also called SGD (Stochastic Gradient Descent) or mini-batch GD.
- What is the effect of using batches?
  - Speeds up convergence.



Figure courtesy of [Ashish Singhal](#)

# Optimization: Batches

- Gradient descent will take the model to the closest minima, not necessarily the global minima.
- By taking batches, we introduce noisiness (randomness) to the loss surface, which may help us avoid local minima.

