

# Intro to Neural Networks

**BA865 – Mohannad Elhamod**

# CNNs

## Convolutional Networks

# A Problem of Scalability

- How many parameters in this network?
- Do we really need to learn all these parameters?

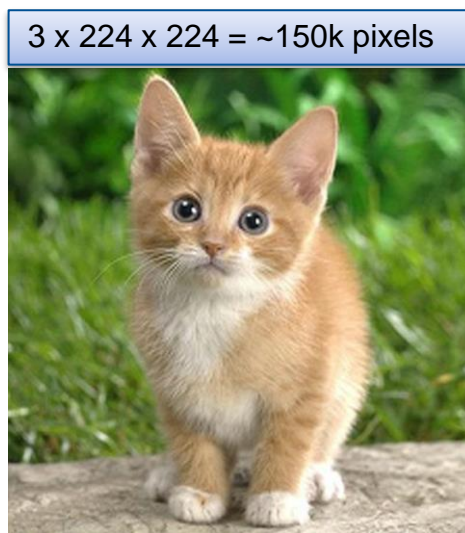


Figure courtesy of Robert Bond

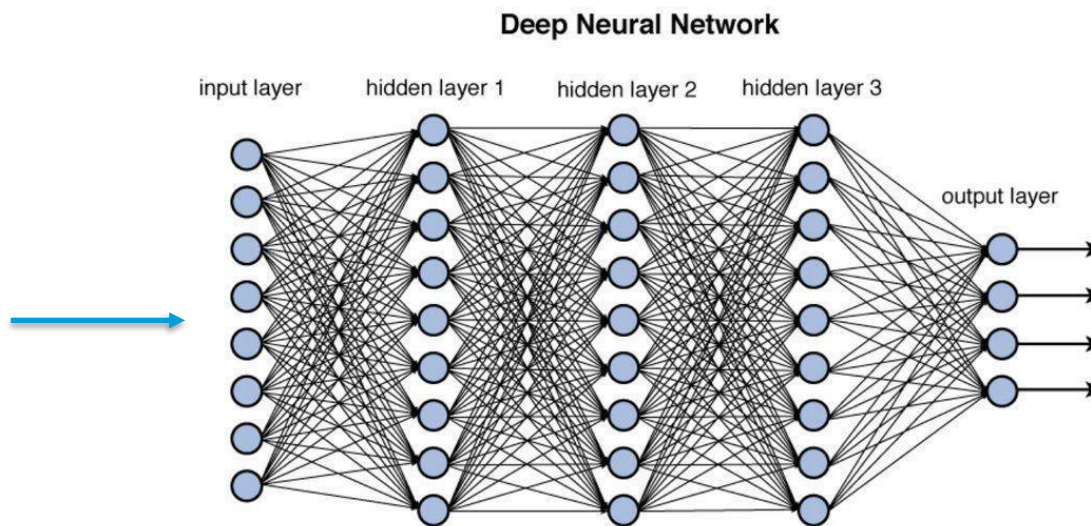


Figure 12.2 Deep network architecture with multiple layers.

Figure courtesy of Ravindra Pamar

# Structure in Images

- Interesting images have:
  - Locality of information.
  - Spatial invariance.



Figure courtesy of Robert Bradi

vs.

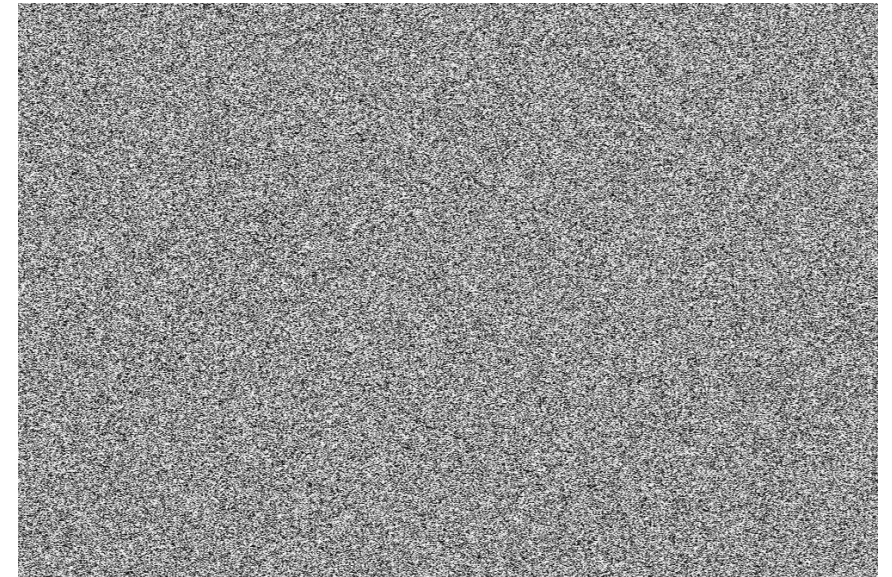


Figure courtesy of Jorge Stolfi

# Convolutional Filters

- Instead of learning a mesh of all possible parameters, let's learn local descriptors (kernels or filters) that can be reused across the image!

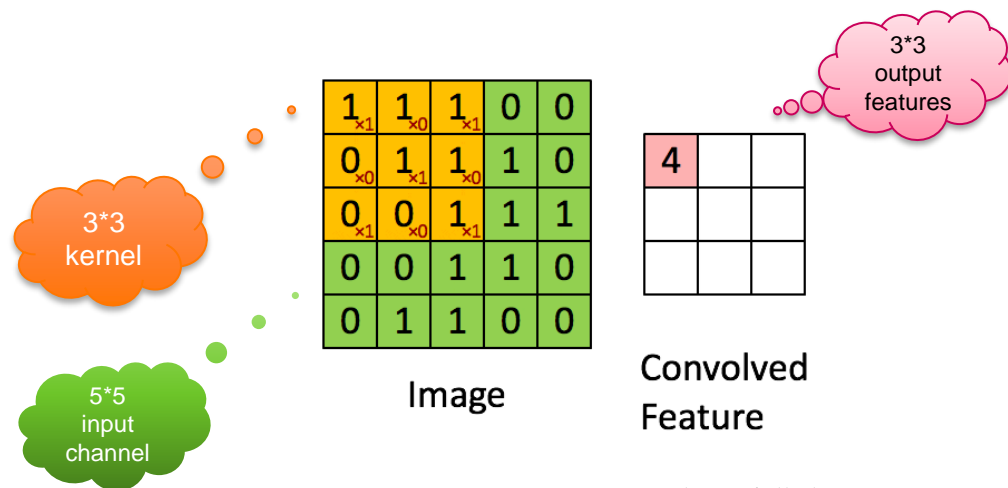


Figure courtesy of Daniel Nouri





Figure courtesy of Thushan Ganegodara

# Mathematically Speaking...

## Convolution Layer in Keras

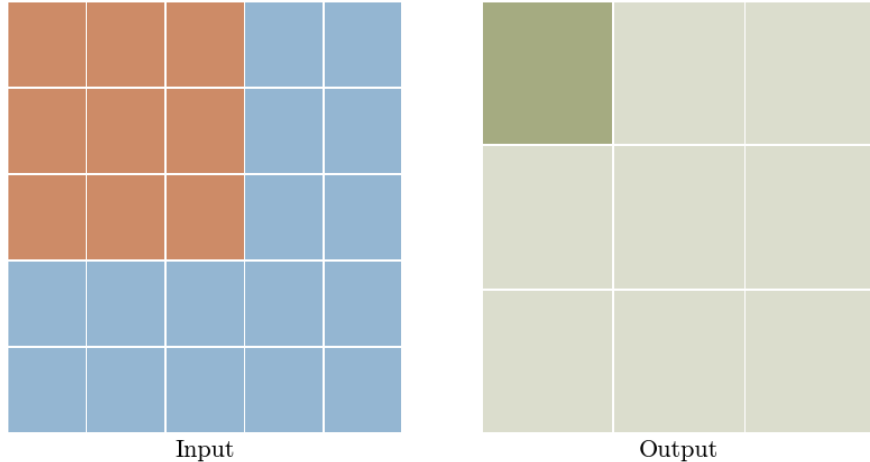




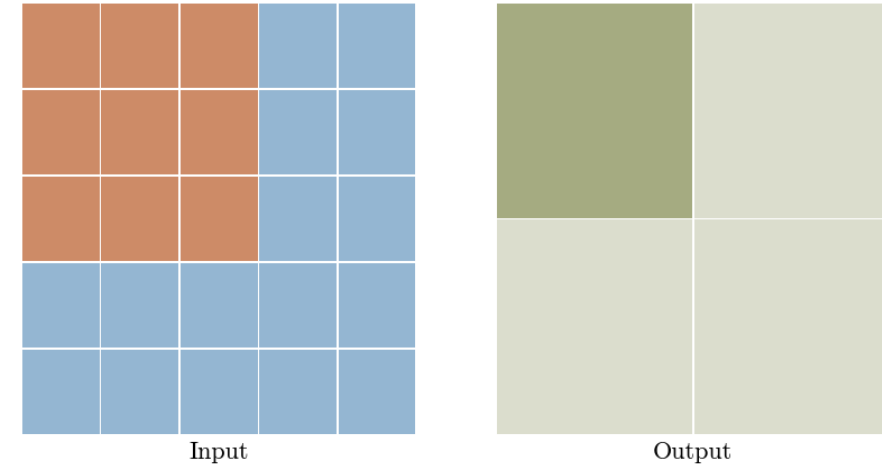
# Stride and Padding

- Stride controls feature field overlap.
- Padding controls the down-sampling.


Type: conv - Stride: 1 Padding: 0



Type: conv - Stride: 2 Padding: 0



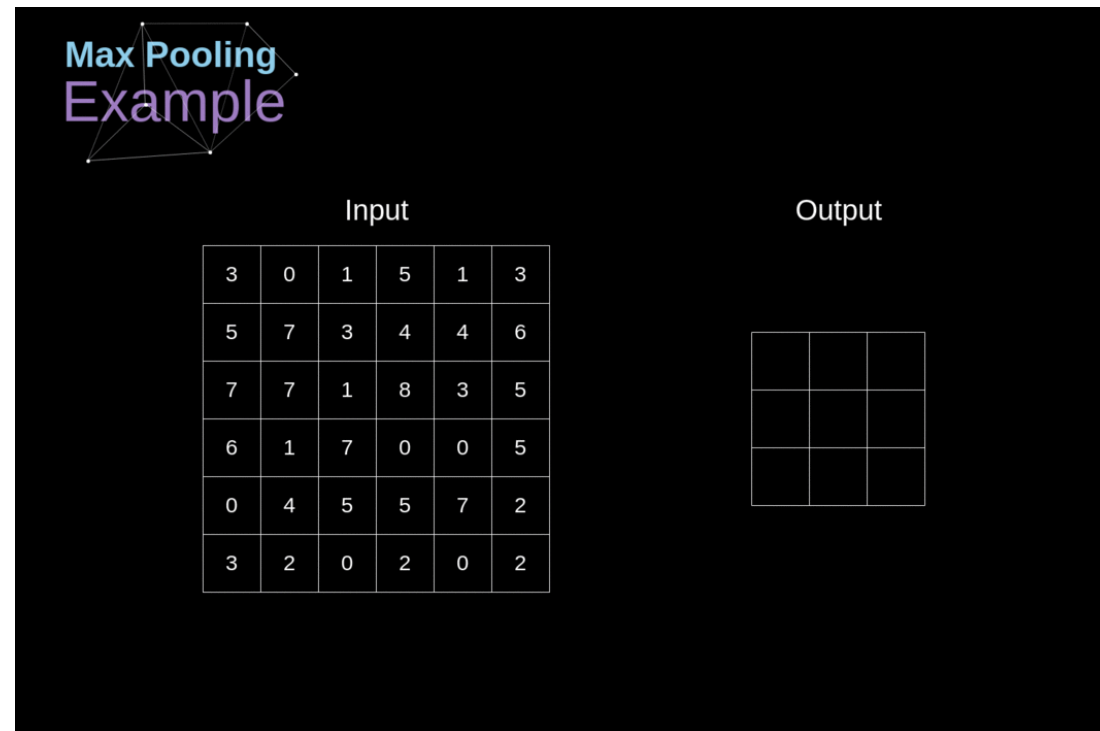
Type: conv - Stride: 1 Padding: 1



[Aqeel Anwar](#)

# Non-Linearity: MaxPooling

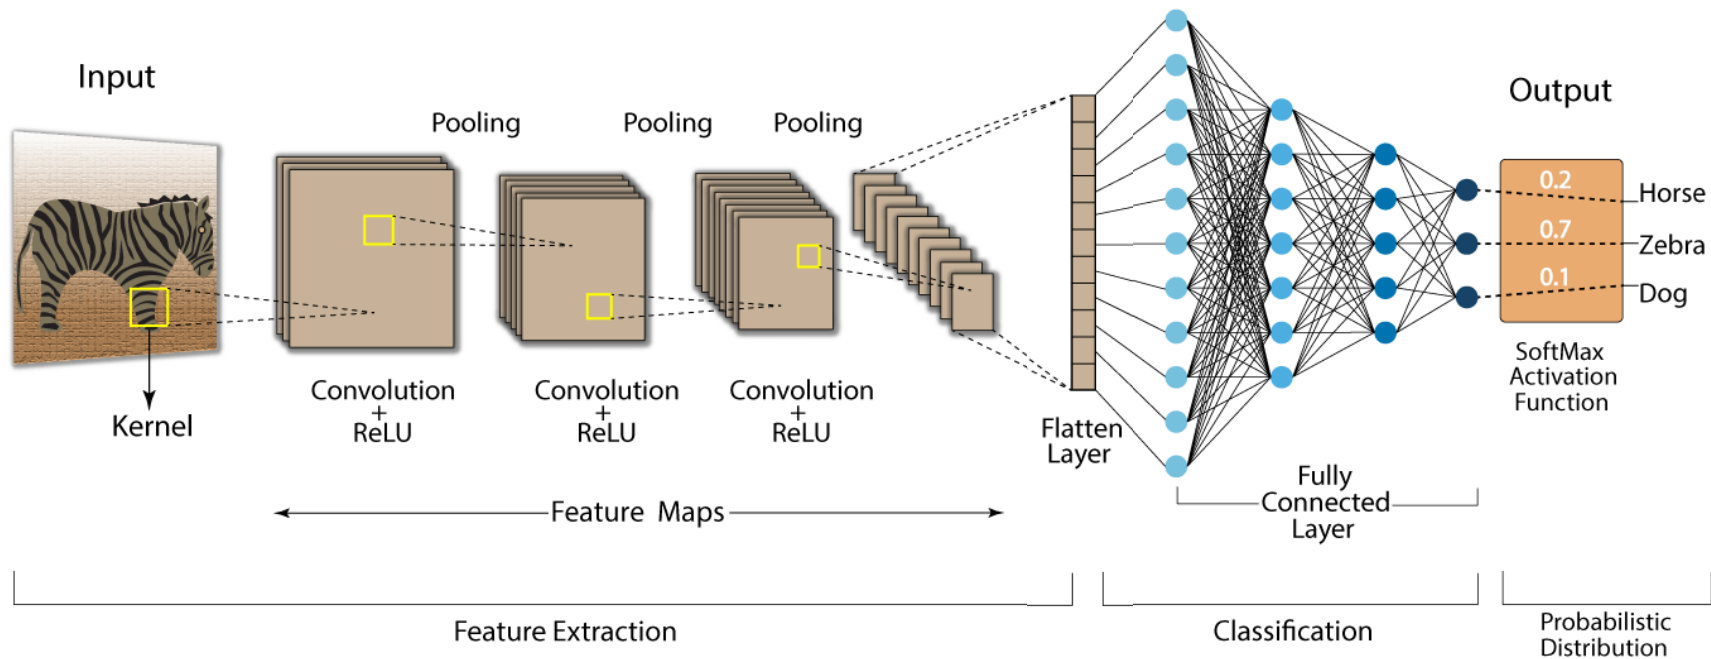
- In addition to being a non-linearity...
  - it helps down-sample the image.
  - It helps summarize information in terms of larger blocks.





# Putting It All Together

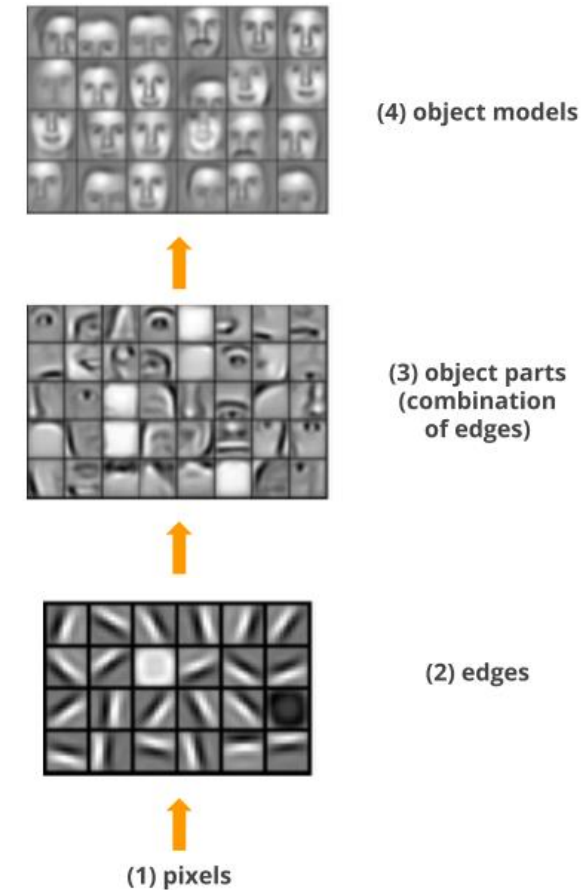
- Deeper layers generally have more kernels that are smaller.



[Afaq Umer](#)

# Learned Features

- Early layers learn low-level features.
  - spots, edges, etc.
- Later layers learn to detect high-level features as a combination of low-level features.
  - Eyes, ears, hair, etc.
- Interpretability is not guaranteed (But there is great research interest...)
- Demo



# Hyper-Parameters

**Continued...**

# Batch Normalization

- Even if input data is properly normalized, the gradient in subsequent layers may vanish or explode.
- For that, you may add “batch normalization” at every layer.

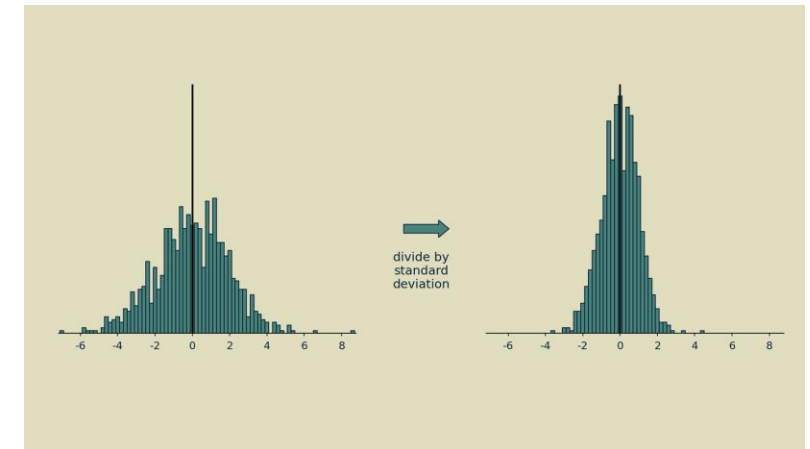
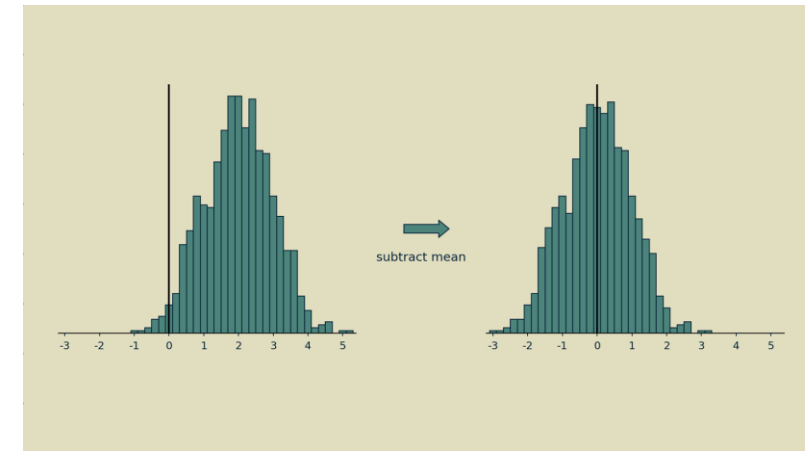


Figure courtesy of Brandon Rohrer

# Learning Rate: Schedulers

- Since larger learning rates may converge faster but smaller ones are more stable, you could adjust the learning rate in phases to get the best of both worlds!
  - This way, you still converge but faster.
- Using a scheduler is a common practice.

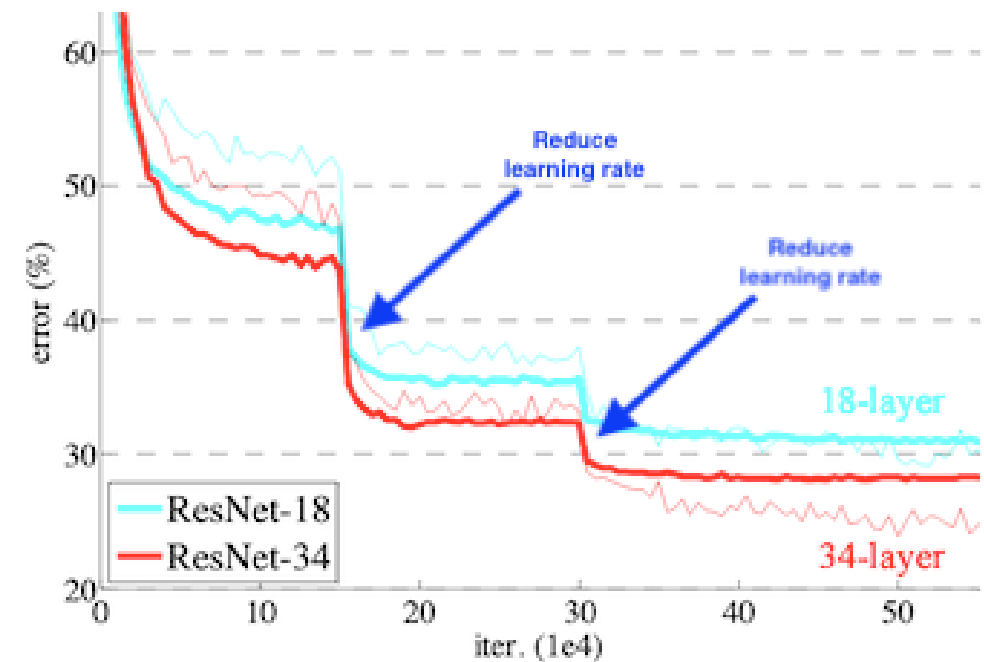


Figure courtesy of [B. D. Hammel](#)

# Hyper- Parameter Tuning

# Be Smart About It

- It is expensive!
  - 1 hyper-parameter with 3 values → 3 experiments
  - 2 hyper-parameter with 3 values each → 9 experiments
  - 3 hyper-parameter with 3 values each → 27 experiments
  - ... exponential growth!



# Be Smart About It

- It is expensive!
- Start with generally accepted wisdom:
  - Start with good initial guesses.
  - Different settings work better for different models/problems (e.g., SGD + momentum for computer vision vs. Adam otherwise)
- Be picky about what to fine-tune.
  - Use early stopping.
  - Learning rate is the most important parameter!

# Hyper-Parameter Tuning Methods

- Generally, use log-scale for numerical hyper-parameters.
- Random and Adaptive searches generally find optimal values faster than grid searches.




Figure courtesy of Liam Li

# Advanced Techniques

# Transfer Learning


- Deeper networks have highly irregular loss surfaces. They are hard to train:
  1. They need relatively large amounts of data.
  2. They need relatively much compute resources to train and tune hyper-parameters.
- We need to somehow start with “*an advantage*”.

# Transfer Learning

- We need to somehow start with “*an advantage*”.
- Large tech companies and research institutes are more capable than individuals in terms of data and compute resources.
  - They can afford to train their models from scratch.
  - Can we capitalize on their pre-trained models?

# Transfer Learning

- A pre-trained model would already have learned useful features for a target problem.
  - For example, we can start with a model (e.g., ResNet) that was pre-trained on a large dataset (e.g., [ImageNet](#): ~1.4M images. 1000 classes. ~3\*469x387 pixels).



# Data Augmentation

- As mentioned earlier, lack of large amounts of data is a problem.
  - Model may overfit (learn “spurious” features).
  - Model may not generalize well to “out-of-distribution” data.



# Data Augmentation

- What is a lion exactly?
- To increase the amount of data and add make sure the learned features are diverse...
  - introduce as much valid variations as possible to the dataset.

