

Intro to Neural Networks

BA865 – Mohannad Elhamod


Auto- Encoders

Compression and Reconstruction

- Given an input (e.g., image), I may want to...
 - compress it and reconstruct it.
 - modify it as it is reconstructed

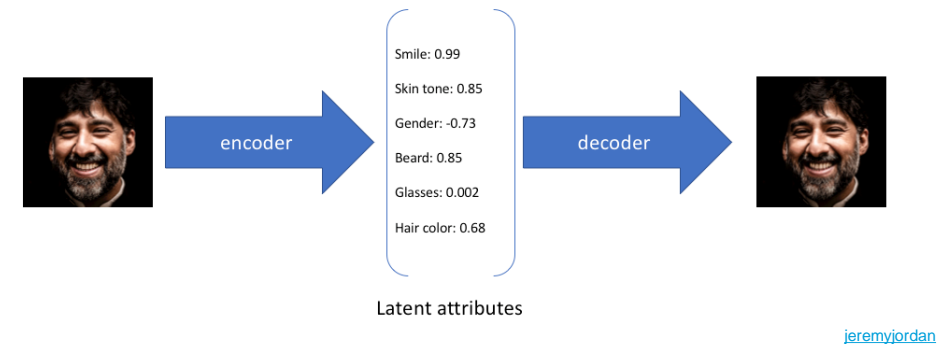

Auto-Encoder

- Given an input, we need to learn a representation (i.e., code, embedding)
- This embedding is the compressed version of the data.
- The embedding should be sufficient to obtain the desired reconstruction.



Auto-Encoder


- Given an input, we need to learn a representation (i.e., code, embedding)
- This embedding is the compressed version of the data
- The embedding should be sufficient to obtain the desired reconstruction.



Auto-Encoder

- It is a non-linear dimensionality reduction method.


Linear vs nonlinear dimensionality reduction



[Herminarto Nugroho](#)

The Error Function

- The error is the “reconstruction loss”
 - The MSE between the input and the output.

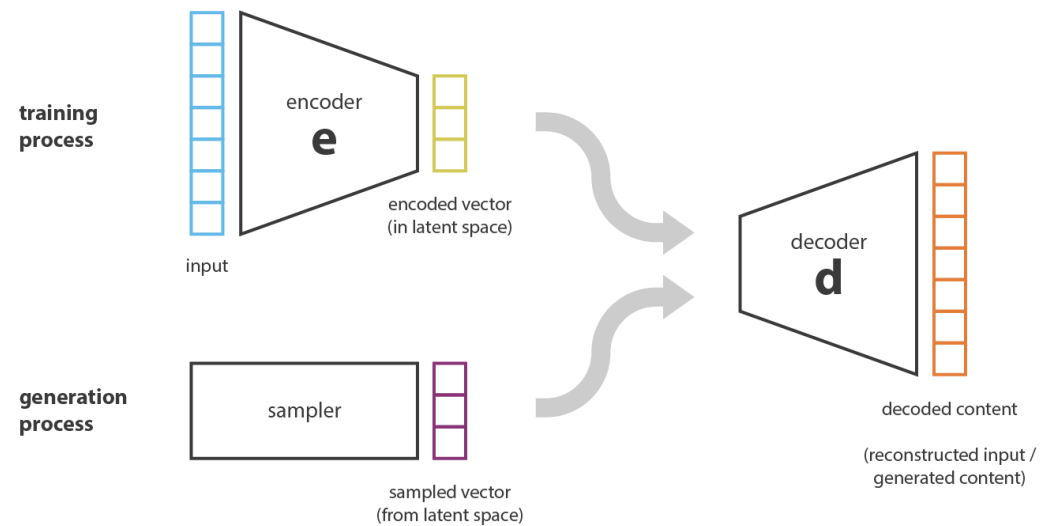


$$\text{loss} = \|x - \hat{x}\|^2 = \|x - d(z)\|^2 = \|x - d(e(x))\|^2$$

[Joseph Rocca](#)

The Auto-Encoder as a Generator


- Once the model is trained, we could use the decoder to generate new content!



[Joseph Rocca](#)

Variational Auto-Encoder

- What if I “make” the embedding to follow a nice Gaussian distribution
- [Demo](#)



$$\text{loss} = ||x - \hat{x}||^2 + \text{KL}[N(\mu_x, \sigma_x), N(0, I)] = ||x - d(z)||^2 + \text{KL}[N(\mu_x, \sigma_x), N(0, I)]$$

[Joseph Rocca](#)

Variational Auto-Encoder

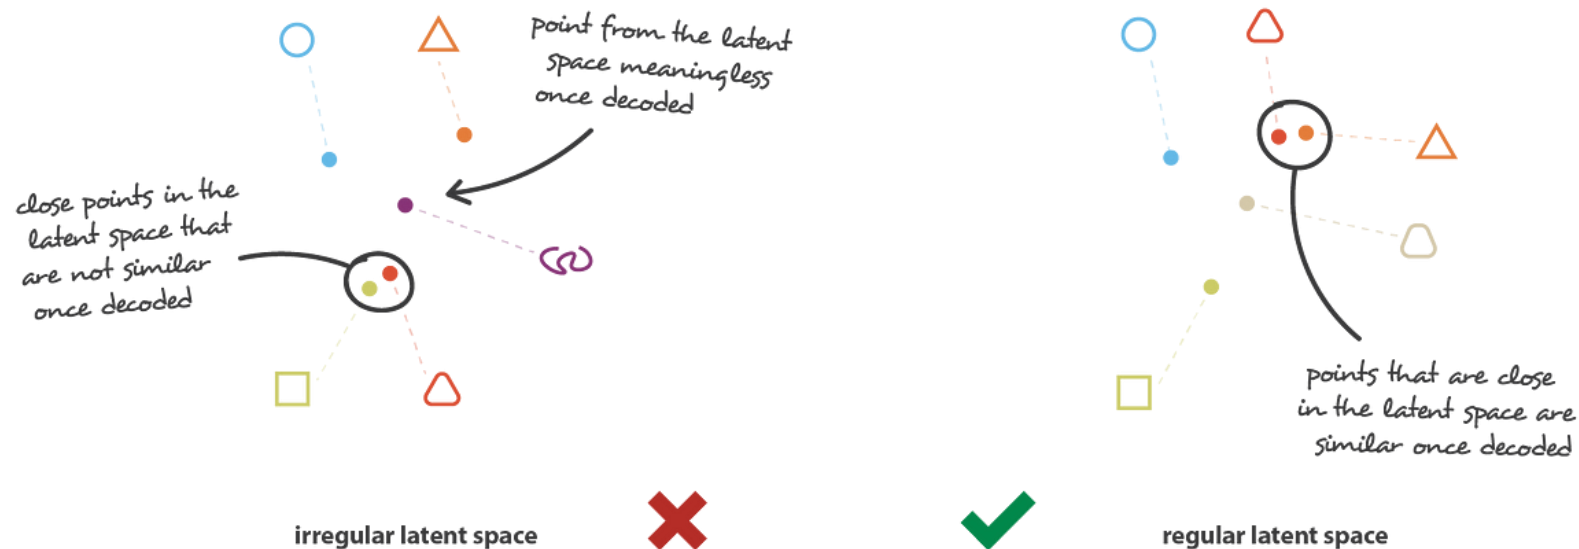
- Consequently, a traversal of the latent space would lead to smoother transitions in the reconstructed data.



[Joseph Rocca](#)



Variational Auto-Encoder

- Consequently, a traversal of the latent space would lead to smoother transitions in the reconstructed data.



[Joseph Rocca](#)

Examples in Computer Vision



Semantic segmentation